



AI Vendor Red Flag Guide

A plain-English buyer's checklist for evaluating AI tools and vendors

Congratulations on completing Module 5. You earned this.

1 10 Red Flags in AI Sales Pitches

What vendors say — and what it actually means.

THEY SAY

"Our AI is powered by proprietary technology."

WHAT IT MEANS

We built a wrapper around an existing AI model (usually OpenAI, Google, or Anthropic) and aren't telling you that. Ask directly: what foundation model powers this? If they won't answer, that's your answer.

THEY SAY

"It learns and improves automatically."

WHAT IT MEANS

Most hosted AI tools don't actually retrain on your data. What they usually mean is your past conversations get stored as context. Ask exactly how personalization works: is it fine-tuning, document retrieval, or just conversation history?

THEY SAY

"99% accuracy."

WHAT IT MEANS

Every AI system makes mistakes. An accuracy claim without context is meaningless. Ask: accurate at what specific task? Measured on what data? Tested by the vendor or an independent party? If they can't answer all three, the number is marketing, not a specification.

THEY SAY

"It replaces [X] employees."

WHAT IT MEANS

Under ideal conditions, with a skilled person operating it, the tool can handle some tasks those employees do. The realistic version: it

handles routine parts, still needs oversight, and your team needs to learn how to use it. Ask what the average implementation looks like, not the best-case story.

THEY SAY

"Set it and forget it."

WHAT IT MEANS

There is no such thing as a "set it and forget it" AI tool. Every deployment requires setup, testing, periodic review, and course corrections. Plan for ongoing maintenance time, not just an initial launch.

THEY SAY

"We integrate with everything."

WHAT IT MEANS

We connect to the 15–20 most popular tools via their standard APIs. Ask for your specific integration to be demonstrated live, with your actual data. Ask what happens when the integration breaks. "Native integration" often means a one-way data push with no error handling.

THEY SAY

"Our AI understands your industry."

WHAT IT MEANS

The underlying model is general-purpose — what's industry-specific are the prompts and curated knowledge base built on top of it. Ask what happens when it produces wrong industry-specific information and whether there's a human review process.

THEY SAY

"No technical setup required — works out of the box."

WHAT IT MEANS

The default demo works immediately. Getting the tool to actually work for your business — connecting your data, configuring workflows, setting guardrails, testing edge cases — takes real work. Plan for at least one to four weeks of setup for any serious deployment.

THEY SAY

"ROI guarantee — 10x your output in 30 days."

WHAT IT MEANS

This is a best-case result, selected from the most successful customer in their portfolio. Ask what the average time-to-value is across all customers, not just the success stories. Ask

what the onboarding timeline looks like step

THEY SAY

"This feature is coming in Q3."

WHAT IT MEANS

That's an aspiration, not a commitment. Roadmap slides are not contracts. If a feature is important to your decision, get a written guarantee with a refund clause, or wait until it ships. Buying based on roadmap promises is how businesses end up paying for tools that don't yet do what they need.

2

12 Questions to Ask Before You Buy

What a good answer looks like — and what a bad one sounds like.

1. What happens to my data?

GOOD ANSWER

"Your data is not used to train the model. Here's our Data Processing Agreement." Check the actual privacy policy for "may be used to improve the service" — that often means your data trains the model unless you opt out.

BAD ANSWER

"We take privacy seriously." — with no specifics.

2. What does it actually cost after the trial — and at 3x my current usage?

GOOD ANSWER

A clear pricing table with usage tiers and overage rates. AI pricing often scales aggressively — a \$99/month tool can reach \$800/month if you actually rely on it.

BAD ANSWER

"We'll figure out a plan that works for you."

3. Can I export all my data if I leave?

GOOD ANSWER

"Yes — here's how, and here's what format it exports in." The best vendors make this easy because they're confident you'll stay.

BAD ANSWER

"Our team can help you with that process." If you can't export independently, you're locked in.

4. What does it do when it doesn't know the answer?

GOOD ANSWER

"Here's how errors surface — the tool shows a confidence score and flags when it's uncertain."

BAD ANSWER

"It rarely gets things wrong." Every AI tool is wrong sometimes. A good vendor explains how errors are handled, not whether errors happen.

5. How often does it need human oversight?

GOOD ANSWER

"For this type of task, you should review outputs before they go out — here's what that workflow looks like."

BAD ANSWER

"The AI handles it automatically — no review needed." Any vendor claiming their AI never needs a second set of eyes is not being honest with you.

6. Is my data used to train the model?

GOOD ANSWER

"No — and here's the signed Data Processing Agreement that confirms it." Check the Terms of Service directly.

BAD ANSWER

"We follow all standard privacy practices." For consumer AI tools, opting out of training data use is often buried in settings — not the default.

7. What does cancellation look like?

GOOD ANSWER

"Cancel anytime, 30 days notice, your data is exportable for 60 days after cancellation."

BAD ANSWER

Unclear, or requires a conversation with a "customer success manager." If the cancellation path is convoluted, it's designed to keep you paying.

8. Can I talk to two or three actual customers in my situation?

GOOD ANSWER

"Absolutely — here are three customers in similar businesses who opted in to speak with prospects." Unscripted conversations reveal what doesn't work.

BAD ANSWER

"We have these case studies on our website." Case studies are curated. If a vendor won't connect you with references, they're hiding something.

9. What's your uptime track record?

GOOD ANSWER

A publicly accessible status page showing 90-plus days of uptime history.

BAD ANSWER

"We have enterprise-grade infrastructure." If the tool becomes part of your daily operations, its downtime is your downtime.

10. What does a failed implementation look like — and how common is it?

GOOD ANSWER

"Here are the most common reasons it doesn't work and what we do differently when we

recognize those patterns." Honest vendors have learned from failures.

BAD ANSWER

Discomfort, deflection, or "it always works when customers follow the process."

11. What does onboarding actually involve, step by step?

GOOD ANSWER

A clear timeline — kickoff, data connection, configuration, testing, go-live — with realistic durations. Ask: who does the work — you, their team, or a third-party partner?

BAD ANSWER

"You'll be up in minutes."

12. What does this cost at 10x my current usage?

GOOD ANSWER

A clear answer with specific numbers at higher volume. The question isn't what it costs today — it's what it costs if the tool actually works and you use it heavily.

BAD ANSWER

"We have plans for businesses of all sizes."
Understand the pricing ceiling before you get there.

3 Do You Actually Need AI?

Answer these five questions honestly before buying or building anything.

QUESTION	YES	NO
Is the task repetitive? You do it often, in roughly the same way each time.	AI might help	You probably just need better organization
Does it involve language or text? Writing, summarizing, answering questions, sorting documents, drafting responses.	AI might help	AI may not be the right tool for this
Is good enough better than	AI might	If accuracy is critical, the cost of

QUESTION	YES	NO
nothing? The task is currently not done at all, or done slowly and inconsistently.	help	AI errors may exceed the benefit
Do you have examples of good output? You can describe or show what a correct result looks like.	AI might help	Without examples, you can't train, test, or verify the AI's work
Can you verify the output? Someone on your team can review AI results and catch mistakes before they cause problems.	AI might help	If no one can verify it, AI errors will be invisible until they cause damage

Answered Yes to most:

AI is worth exploring. Start with a low-commitment trial on one specific task before rolling it out broadly.

Answered No to most:

You may not have an AI problem. You may have a process problem. Better workflows, clearer responsibilities, or simpler tools might solve it faster and cheaper.

One more question first:

Is this task being done poorly because of missing technology — or because of missing time, unclear ownership, or inconsistent processes? AI layered on top of a broken process produces broken results faster. Fix the process first.

4 Pricing Models Decoded

A quick reference for the most common AI pricing structures.

Per Seat / Per User

You pay a fixed monthly fee for each person who has access. Example: \$50 per user per month.

WORKS WELL WHEN

Adoption is high and all users actively use the tool

WATCH OUT FOR

Paying for inactive seats — half your team rarely opens it, half your spend is wasted

Per API Call / Per Token

You pay based on what you actually use.
Example: \$0.002 per 1,000 words processed.

WORKS WELL WHEN

Usage is variable or hard to predict upfront

WATCH OUT FOR

Surprise bills — one large batch job can generate a very large bill overnight. Always set a spending cap first.

Flat Monthly Fee

One price for access up to a defined limit (or "unlimited," with conditions). Example: \$299 per month.

WORKS WELL WHEN

High, consistent usage and predictable cost is the priority

WATCH OUT FOR

What "unlimited" actually means — most plans have a soft limit where throttling or overage charges kick in

Usage-Based / Metered (Tiered)

A flat rate at each usage tier, with step-up pricing as you grow. Example: \$49/month up to 5,000 messages, \$299/month for 5,001–25,000 messages.

WORKS WELL WHEN

Starting small as a low-commitment entry

WATCH OUT FOR

Steep jumps between tiers — your cost can

point

jump 5x with one month of growth. Project at 3x and 10x usage before committing.

"Free Forever"

The core product is free. The features you actually need are paid add-ons. By the time you've added everything you need, the "free" tool often costs more than a competing product with transparent flat-rate pricing.

WORKS WELL WHEN

Your needs are genuinely covered by the free tier — rare

WATCH OUT FOR

Factor all add-ons into your comparison, not just the headline price. The demo is free. Useful features are not.

The Cost the Pricing Page Never Shows

Whatever the pricing page says, the real first-year cost is higher. A \$99/month tool that takes 40 hours of setup has a real first-year cost thousands of dollars higher than the subscription alone.

Time to set it up and configure it for your business

Time to connect it to your existing tools

Time to train your team

Time spent reviewing and correcting AI outputs

The cost of switching to something else if it doesn't work out

The best AI vendors welcome hard questions.

They've built real products for real customers, they know where the edges are, and they're confident enough in what they've built to be honest about what it doesn't do. If asking any of these questions makes a vendor uncomfortable, defensive, or evasive — that tells you everything you need to know.

A free resource from **Epiphany Dynamics**



"The best vendors welcome hard questions."